

Final-Panel Comments Matrix

One row is retained for each of the 15 reported panel comments and emphasis notes. The matrix separates manuscript corrections already completed from empirical work that still requires new data, annotation, or operational testing.

Status: Addressed = present in the authoritative Markdown. Addressed; validation pending = wording is corrected, but the requested empirical extension remains future work. Priority defense point = written correction is complete and should be emphasized during questioning.

No.	Panel comment	Required interpretation	Action taken / evidence	Manuscript location	Status and remaining work
1	The model has limited ability to detect predator conversations before a sufficient behavioral trajectory has developed.	Treat early evidence scarcity as a measured performance issue, not as proof that the full model is invalid.	Chapter IV now identifies the principal observed failure pattern: all four primary-LSTM false negatives had 2, 11, 13, or 26 turns, below the 36-turn first quartile of correctly detected positives. Chapter V records early-trajectory performance as a focused next problem.	Chapter IV, Section 4.6; Chapter V, Section 5.1, Finding 6	Addressed; validation pending. A future prefix-based study should measure recall and time-to-detection at fixed turn counts using expert onset annotations.
2	The study can retain the model, but the reasoning must be careful.	Explain why the trajectory model remains valuable while avoiding an unsupported early-onset claim.	The LSTM is presented confidently as the strongest primary conversation-prioritization method, supported by the matched seven-feature comparison. It is also positioned as a cumulative trajectory component that can be paired with immediate content safeguards.	Chapter IV, Sections 4.3, 4.6, and 4.7; Chapter V, Section 5.2	Addressed. Do not describe the current model as a validated grooming-onset detector.
3	Evaluate latency for "real-time" chat users; otherwise remove or rephrase the claim.	Separate offline sequential capability from demonstrated production real-time performance.	The manuscript now calls the current process controlled offline sequential replay. It reserves "production real-time" claims until end-to-end latency, throughput, memory use, and representative-traffic tests are completed.	Chapter I, Section 1.4.2; Chapter III, Section 3.1; Chapter IV, Section 4.7; Chapter V, Recommendation 3	Addressed; validation pending. Run a frozen-model latency benchmark and report median, p95, and p99 per-turn latency plus throughput before claiming real-time deployment.
4	PAN12 is not representative of Gen Z or present-day communication; determine how grooming behavior appears today.	Treat dataset age as an external-validity question while preserving the completed benchmark result.	PAN-2012 is described as the historical benchmark selected to evaluate the architecture. The paper calls for external validation on contemporary communication, current slang and obfuscation, Filipino/Taglish data, and representative platform environments.	Chapter I, Section 1.4.2; Chapter V, Recommendations 1 and 4	Addressed; validation pending. Contemporary independently labeled data must be acquired and evaluated as a separate external-validation study.
5	When presenting at a conference, explain what the research actually does.	State one precise system purpose and evaluation target.	The study now states that the module analyzes chronological contextual and behavioral evidence and assigns a conversation-level priority for human moderator review. It does not issue autonomous sanctions.	Chapter I, Sections 1.2 and 1.4.1; Chapter III, Section 3.1; Chapter IV, Section 4.7; Chapter V, Section 5.2	Addressed. Conference answer: "The system prioritizes conversations for moderator review by learning how contextual evidence develops across turns."
6	Define whether the system detects users who are predators or detects grooming behavior.	Distinguish the training supervision from the evaluated output.	Layer 1 is explicitly trained from official predator-author membership and produces a predator-author proxy score. Layer 2 uses the score trajectory to classify the conversation-level endpoint. The output is not a user-level offender verdict.	Chapter I, Sections 1.2 and 1.6; Chapter III, Sections 3.1, 3.3.1, and 3.4.1	Addressed. The system prioritizes conversations; author membership is the supervision source, not the final enforcement decision.
7	Clarify whether the study is looking at grooming behavior or only at conversations.	Name the completed unit of analysis and avoid switching units in the results.	The operative output and primary evaluation unit are consistently defined as the conversation. Turn-level scores are intermediate sequence signals rather than validated grooming-message labels.	Chapter I, Sections 1.2 and 1.6; Chapter III, Sections 3.3.1 and 3.5.2; Chapter IV, Sections 4.1 and 4.6	Addressed. All reported confusion counts and performance metrics use conversations as the unit.
8	Detecting grooming behavior and identifying conversations that may lead to grooming are different tasks.	Prevent current conversation prioritization from being described as direct behavioral or future-harm prediction.	The manuscript distinguishes conversation prioritization from message-level behavior labeling, intent determination, and onset prediction. Human review remains responsible for interpreting content and deciding action.	Chapter I, Sections 1.2 and 1.4.2; Chapter III, Sections 3.5.2 and 3.6	Addressed. Avoid phrases such as "predicts that grooming will occur" unless a later prospective study directly tests that claim.

No.	Panel comment	Required interpretation	Action taken / evidence	Manuscript location	Status and remaining work
9	The presentation appears to describe three targets: behavior, conversation, and grooming onset.	Select one validated target and classify the other two as separate research extensions.	The completed target is conversation-level prioritization. Behavioral trajectory features are model inputs, not behavioral ground-truth labels. Grooming onset is a future annotation and evaluation task.	Chapter I, Section 1.2; Chapter III, Sections 3.3.4 and 3.5.2; Chapter V, Recommendation 1	Addressed. One current endpoint, two clearly separated extensions.
10	Clarify the methodology for complete grooming-onset annotation and how onset ground truth will be verified.	Do not imply that onset ground truth exists when it was never independently annotated.	Prefix scores are now described only as showing how priority changes as evidence accumulates. The recommendations require expert-reviewed first-onset annotations before evaluating time-to-detection.	Chapter III, Section 3.5.2; Chapter IV, Section 4.6; Chapter V, Recommendation 1	Addressed; validation pending. A future onset protocol needs trained reviewers, an operational definition, independent annotation, adjudication, and inter-rater agreement.
11	Motivate the trajectory features using OGDM.	Show a defensible theory-to-computation mapping without pretending the features are OGDM stage labels.	Chapter II now maps OGDM-informed conversational properties to the seven feature families: accumulation and persistence, escalation and retreat, semantic movement, and interaction structure. Chapter IV interprets the same mapping in the results.	Chapter II, Section 2.3, OGDM mapping table; Chapter IV, Section 4.5	Addressed. The features are theory-informed operational indicators, while the LSTM learns their chronological interaction.
12	Provide more evidence and explanation for author partitioning.	Explain both the procedure and the leakage it prevents.	Conversations are graph vertices; shared authors connect them; every connected component is assigned wholly to one partition. This creates zero author and component overlap and prevents recurring participant language from leaking into development and evaluation. Component-grouped bootstrap resampling preserves the same dependency structure.	Chapter III, Sections 3.3.3 and 3.5.2; Chapter IV, Sections 4.1, 4.3, and 4.4	Addressed. Final-test performance represents transfer to unseen author-connected groups within the selected benchmark.
13	F0.5 gives more weight to precision, yet the study emphasizes recall. Explain the choice and be prepared to defend it.	Separate the operating-threshold objective from the safety outcomes used to answer the research objective.	F0.5 selects a precision-oriented operating point for a finite human-review queue, where excessive false positives consume reviewer capacity. Recall, false negatives, PR-AUC, and the full confusion matrix remain separately reported; the threshold is not described as recall-optimized.	Chapter III, Section 3.5.2; Chapter IV, Section 4.7; Chapter V, Section 5.2, Specific Objective 3	Addressed; priority defense point. The model achieved 0.8511 precision, 0.9091 recall, F0.5 of 0.8621, seven false positives, and four false negatives.
14	Further findings could make the work a stronger child-safety application; explain the results more fully.	Connect the empirical improvement to a realistic moderation workflow without claiming proven harm prevention.	Chapter IV quantifies the LSTM's gains over keyword, maximum Layer 1, and weighted scoring, then explains conversation prioritization for sequential human review. Chapter V directly answers every unchanged research objective and specifies the next deployment studies.	Chapter IV, Sections 4.2-4.8; Chapter V, Sections 5.1-5.3	Addressed; validation pending. Future work should evaluate moderator effectiveness, calibration, latency, contemporary data, privacy, fairness, and human escalation outcomes.
15	The panel is especially likely to focus on F0.5.	Prepare a short, numerically grounded explanation rather than treating the metric as interchangeable with recall.	The methodology and conclusions now state why F0.5 is used and how it coexists with the recall-focused objective. The accepted result table supplies the exact figures and comparator differences.	Chapter III, Section 3.5.2; Chapter IV, Tables 4.1 and 4.2; Chapter V, Section 5.2	Addressed; priority defense point. Ready answer: "F0.5 chose an actionable precision-oriented threshold, while recall and false-negative counts independently tested whether the model reduced missed positive conversations."

Consolidated Status

- Manuscript corrections completed: Comments 2, 5-9, 11-13, and the written portions of 1, 3, 4, 10, 14, and 15.
- New empirical work still required: early-prefix/onset evaluation, production latency and throughput testing, contemporary external-data validation, and moderator-impact evaluation.
- No change permitted or required: the adviser-locked research questions and objectives, frozen model artifacts, selected thresholds, held-out results, and consumed final-test ledger.
- Current synchronization boundary: these panel revisions are present in Markdown only. The Word manuscripts and presentation have not yet been updated from this revision set.

Core Defense Statement

The research developed and evaluated a complete two-layer moderation module that prioritizes conversations for human review by modeling how contextual evidence develops across turns. On the author-disjoint held-out test, the primary seven-feature LSTM achieved 0.9153 PR-AUC, 0.8621 F0.5, 0.8511 precision, and 0.9091 recall, and it significantly improved PR-AUC and F0.5 over the weighted scorer using the same seven inputs. The results establish the value of learned recurrent trajectory modeling for conversation prioritization, while early-onset detection, contemporary-data transfer, and production latency remain clearly defined next evaluation stages.